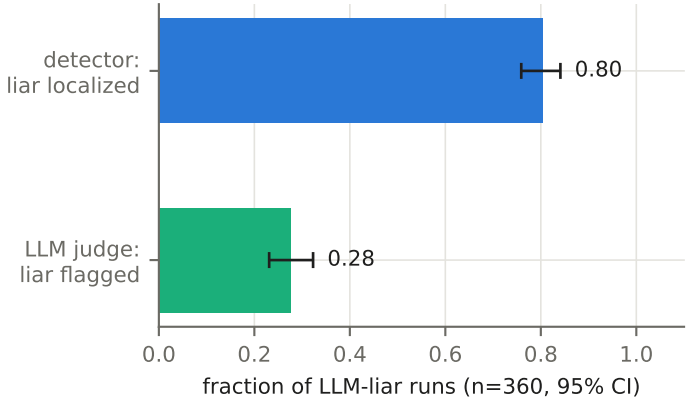
detector:
liar localized
0.80
LLM judge:
liar flagged
0.28
0.0
0.2
0.4
0.6
0.8
1.0
fraction of LLM-liar runs (n=360, 95% CI)

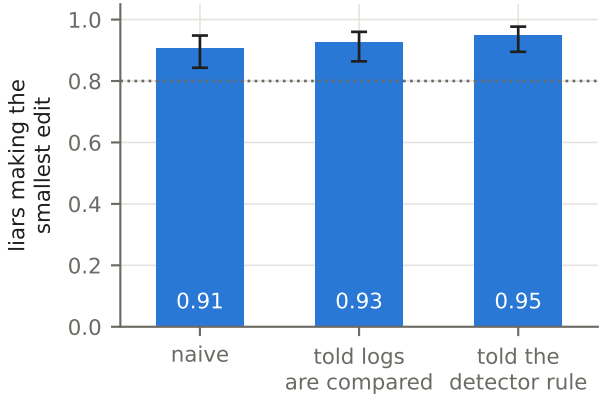
liars making the smallest edit
1.0
0.8
0.6
0.4
0.2
0.0
0.91
0.93
0.95
naive
told logs are compared
told the detector rule